

Teclov_Visualising_Time_Series_Data

March 24, 2020

1 Teclov : Visualising Time Series Data

In the section, we will explore ways to visualise data gathered over time. We will: - Plot simple time series plots - Derive variables such as month and year and use them for richer visualisations

```
[1]: # loading libraries and reading the data

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# set seaborn theme if you prefer
sns.set(style="white")

# read data
market_df = pd.read_csv("./global_sales_data/market_fact.csv")
customer_df = pd.read_csv("./global_sales_data/cust_dimen.csv")
product_df = pd.read_csv("./global_sales_data/prod_dimen.csv")
shipping_df = pd.read_csv("./global_sales_data/shipping_dimen.csv")
orders_df = pd.read_csv("./global_sales_data/orders_dimen.csv")
```

1.1 Visualising Simple Time Series Data

Let's say you want to visualise numeric variables such as Sales, Profit, Shipping_Cost etc. over time.

```
[2]: market_df.head()
```

```
[2]:
```

	Ord_id	Prod_id	Ship_id	Cust_id	Sales	Discount	Order_Quantity	\
0	Ord_5446	Prod_16	SHP_7609	Cust_1818	136.81	0.01	23	
1	Ord_5406	Prod_13	SHP_7549	Cust_1818	42.27	0.01	13	
2	Ord_5446	Prod_4	SHP_7610	Cust_1818	4701.69	0.00	26	
3	Ord_5456	Prod_6	SHP_7625	Cust_1818	2337.89	0.09	43	
4	Ord_5485	Prod_17	SHP_7664	Cust_1818	4233.15	0.08	35	

	Profit	Shipping_Cost	Product_Base_Margin
0	-30.51	3.60	0.56

1	4.56	0.93	0.54
2	1148.90	2.50	0.59
3	729.34	14.30	0.37
4	1219.87	26.30	0.38

Since the `Order_Date` variable is in the orders dataframe, let's merge it.

```
[3]: # merging with the Orders data to get the Date column
df = pd.merge(market_df, orders_df, how='inner', on='Ord_id')
df.head()
```

```
[3]:      Ord_id  Prod_id  Ship_id  Cust_id  Sales  Discount  Order_Quantity  \
0  Ord_5446  Prod_16  SHP_7609  Cust_1818  136.81      0.01             23
1  Ord_5446  Prod_4   SHP_7610  Cust_1818  4701.69      0.00             26
2  Ord_5446  Prod_6   SHP_7608  Cust_1818  164.02      0.03             23
3  Ord_5406  Prod_13  SHP_7549  Cust_1818   42.27      0.01             13
4  Ord_5456  Prod_6   SHP_7625  Cust_1818  2337.89      0.09             43
```

	Profit	Shipping_Cost	Product_Base_Margin	Order_ID	Order_Date	\
0	-30.51	3.60	0.56	36262	27-07-2010	
1	1148.90	2.50	0.59	36262	27-07-2010	
2	-47.64	6.15	0.37	36262	27-07-2010	
3	4.56	0.93	0.54	20513	07-07-2009	
4	729.34	14.30	0.37	39682	09-11-2010	

	Order_Priority
0	NOT SPECIFIED
1	NOT SPECIFIED
2	NOT SPECIFIED
3	HIGH
4	MEDIUM

```
[4]: # Now we have the Order_Date in the df
# It is stored as a string (object) currently
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 8399 entries, 0 to 8398
Data columns (total 13 columns):
Ord_id      8399 non-null object
Prod_id     8399 non-null object
Ship_id     8399 non-null object
Cust_id     8399 non-null object
Sales       8399 non-null float64
Discount    8399 non-null float64
Order_Quantity 8399 non-null int64
Profit      8399 non-null float64
Shipping_Cost 8399 non-null float64
```

```

Product_Base_Margin    8336 non-null float64
Order_ID               8399 non-null int64
Order_Date             8399 non-null object
Order_Priority         8399 non-null object
dtypes: float64(5), int64(2), object(6)
memory usage: 918.6+ KB

```

Since `Order_Date` is a string, we need to convert it into a `datetime` object. You can do that using `pd.to_datetime()`.

```

[5]: # Convert Order_Date to datetime type
df['Order_Date'] = pd.to_datetime(df['Order_Date'])

# Order_Date is now datetime type
df.info()

```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 8399 entries, 0 to 8398
Data columns (total 13 columns):
Ord_id          8399 non-null object
Prod_id         8399 non-null object
Ship_id         8399 non-null object
Cust_id         8399 non-null object
Sales           8399 non-null float64
Discount        8399 non-null float64
Order_Quantity  8399 non-null int64
Profit          8399 non-null float64
Shipping_Cost   8399 non-null float64
Product_Base_Margin 8336 non-null float64
Order_ID        8399 non-null int64
Order_Date      8399 non-null datetime64[ns]
Order_Priority  8399 non-null object
dtypes: datetime64[ns](1), float64(5), int64(2), object(5)
memory usage: 918.6+ KB

```

Now, since on each day, multiple orders were placed, we need to aggregate `Sales` using a metric such as mean, median etc., and then create a time series plot.

We will group by `Order_Date` and compute the sum of `Sales` on each day.

```

[6]: # aggregating total sales on each day
time_df = df.groupby('Order_Date')['Sales'].sum()
print(time_df.head())

print(type(time_df))

```

```

Order_Date
2009-01-01    1052.8400
2009-01-02    5031.9000
2009-01-03    7288.1375

```

```
2009-01-04    6188.4245
2009-01-05    2583.3300
Name: Sales, dtype: float64
<class 'pandas.core.series.Series'>
```

We can now create a time-series plot using `sns.tsplot()`.


```
[7]: # time series plot

# figure size
plt.figure(figsize=(16, 8))

# tsplot
sns.tsplot(data=time_df)
plt.show()
```

/Users/shivank/anaconda/lib/python3.5/site-packages/seaborn/timeseries.py:183:
UserWarning: The tsplot function is deprecated and will be removed or replaced
(in a substantially altered version) in a future release.

```
warnings.warn(msg, UserWarning)
```



1.1.1 Using Derived Date Metrics for Visualisation

It is often helpful to use derived variables from date such as month and year and using them to identify hidden patterns.

```
[8]: # extracting month and year from date

# extract month
df['month'] = df['Order_Date'].dt.month
```

```
# extract year
df['year'] = df['Order_Date'].dt.year

df.head()
```

```
[8]:
```

	Ord_id	Prod_id	Ship_id	Cust_id	Sales	Discount	Order_Quantity	\
0	Ord_5446	Prod_16	SHP_7609	Cust_1818	136.81	0.01	23	
1	Ord_5446	Prod_4	SHP_7610	Cust_1818	4701.69	0.00	26	
2	Ord_5446	Prod_6	SHP_7608	Cust_1818	164.02	0.03	23	
3	Ord_5406	Prod_13	SHP_7549	Cust_1818	42.27	0.01	13	
4	Ord_5456	Prod_6	SHP_7625	Cust_1818	2337.89	0.09	43	

	Profit	Shipping_Cost	Product_Base_Margin	Order_ID	Order_Date	\
0	-30.51	3.60	0.56	36262	2010-07-27	
1	1148.90	2.50	0.59	36262	2010-07-27	
2	-47.64	6.15	0.37	36262	2010-07-27	
3	4.56	0.93	0.54	20513	2009-07-07	
4	729.34	14.30	0.37	39682	2010-09-11	

	Order_Priority	month	year
0	NOT SPECIFIED	7	2010
1	NOT SPECIFIED	7	2010
2	NOT SPECIFIED	7	2010
3	HIGH	7	2009
4	MEDIUM	9	2010

Now you can plot the average sales across years and months.

```
[9]: # grouping by year and month
df_time = df.groupby(["year", "month"]).Sales.mean()
df_time.head()
```

```
[9]:
```


year	month	
2009	1	1898.475090
	2	2116.510723
	3	2258.661599
	4	2374.155868
	5	1922.317055

Name: Sales, dtype: float64

```
[10]: plt.figure(figsize=(8, 6))
# time series plot
sns.tsplot(df_time)
plt.xlabel("Time")
plt.ylabel("Sales")
plt.show()
```

/Users/shivank/anaconda/lib/python3.5/site-packages/seaborn/timeseries.py:183:
 UserWarning: The tsplot function is deprecated and will be removed or replaced
 (in a substantially altered version) in a future release.

warnings.warn(msg, UserWarning)



There is another way to visualise numeric variables, such as `Sales`, across the year and month. We can pivot the `month` column to create a wide-format dataframe, and then plot a heatmap.

```
[11]: # Pivoting the data using 'month'
year_month = pd.pivot_table(df, values='Sales', index='year', columns='month',
                             ↪aggfunc='mean')
year_month.head()
```

```
[11]: month      1      2      3      4      5  \
year
2009  1898.475090  2116.510723  2258.661599  2374.155868  1922.317055
2010  1549.664361  1591.532297  1708.934944  1461.935539  1440.393540
2011  1321.671562  1733.378070  1822.860614  2062.716921  1822.033936
2012  2220.831551  2438.166961  1578.284028  1325.253694  1865.744629

month      6      7      8      9     10  \
year
```


2009	1548.093259	1909.084469	1827.095921	1717.295494	2203.961070
2010	1695.397085	1272.653408	1277.917503	2216.599136	1927.250907
2011	1655.599644	1515.576260	1471.700741	1430.616748	1851.404564
2012	2192.228263	1824.935942	1090.950643	1628.963780	1636.951775

month	11	12
year		
2009	1502.192743	2209.047484
2010	1806.324648	1869.720956
2011	2161.479056	1672.187881
2012	1935.183039	1747.047528

You can now create a heatmap using `sns.heatmap()`.

```
[12]: # figure size
plt.figure(figsize=(12, 8))

# heatmap with a color map of choice
sns.heatmap(year_month, cmap="YlGnBu")
plt.show()
```



1.1.2 Additional Reading on Time Series Plots and Heatmaps

1. Seaborn heatmaps (documentation)